

Problem Set 1

Applied Stats II

Due: February 11, 2026

Instructions

- Please show your work! You may lose points by simply writing in the answer. If the problem requires you to execute commands in R, please include the code you used to get your answers. Please also include the .R file that contains your code. If you are not sure if work needs to be shown for a particular problem, please ask.
- Your homework should be submitted electronically on GitHub in .pdf form.
- This problem set is due before 23:59 on Wednesday February 11, 2026. No late assignments will be accepted.

Question 1

The Kolmogorov-Smirnov test uses cumulative distribution statistics test the similarity of the empirical distribution of some observed data and a specified PDF, and serves as a goodness of fit test. The test statistic is created by:

$$D = \max_{i=1:n} \left\{ \frac{i}{n} - F_{(i)}, F_{(i)} - \frac{i-1}{n} \right\}$$

where F is the theoretical cumulative distribution of the distribution being tested and $F_{(i)}$ is the i th ordered value. Intuitively, the statistic takes the largest absolute difference between the two distribution functions across all x values. Large values indicate dissimilarity and the rejection of the hypothesis that the empirical distribution matches the queried theoretical distribution. The p-value is calculated from the Kolmogorov- Smirnoff CDF:

$$p(D \leq d) = \frac{\sqrt{2\pi}}{d} \sum_{k=1}^{\infty} e^{-(2k-1)^2 \pi^2 / (8d^2)}$$

which generally requires approximation methods (see Marsaglia, Tsang, and Wang 2003). This so-called non-parametric test (this label comes from the fact that the distribution of the test statistic does not depend on the distribution of the data being tested) performs

poorly in small samples, but works well in a simulation environment. Write an R function that implements this test where the reference distribution is normal. Using R generate 1,000 Cauchy random variables (`rcauchy(1000, location = 0, scale = 1)`) and perform the test (remember, use the same seed, something like `set.seed(123)`, whenever you're generating your own data).

As a hint, you can create the empirical distribution and theoretical CDF using this code:

```

1 # create empirical distribution of observed data
2 ECDF <- ecdf(data)
3 empiricalCDF <- ECDF(data)
4 # generate test statistic
5 D <- max(abs(empiricalCDF - pnorm(data)))

```

```

1 rm(data)
2 set.seed(123)
3 data <- rcauchy(1000, location = 0, scale = 1)
4 ks_normal_test <- function(data) {
5   n <- length(data)
6   ECDF <- ecdf(data)
7   empiricalCDF <- ECDF(data)
8   D <- max(abs(empiricalCDF - pnorm(data)))
9   k <- 1:100
10  p_value <- 2 * sum((-1)^(k - 1) * exp(-2 * k^2 * n * D^2))
11  p_value <- max(min(p_value, 1), 0)
12  return(list(
13    D = D,
14    p.value = p_value))
15 set.seed(123)
16 data <- rcauchy(1000, location = 0, scale = 1)
17 result <- ks_normal_test(data)
18 result

```

The code first ensures that `data` refers to a numeric vector rather than a function, then creates the empirical cumulative distribution function `ECDF` for the observed data. A custom function, `ks normal test`, is defined to perform a Kolmogorov-Smirnov test against a normal distribution. Within this function, the KS statistic D is calculated as the maximum absolute difference between the empirical CDF and the theoretical normal CDF, and the corresponding p-value is approximated using the Kolmogorov distribution. The test is then applied to 1,000 randomly generated Cauchy values. The resulting large D and small p-value indicate a substantial deviation from normality, leading to rejection of the null hypothesis that the data are normally distributed.

- $D = 0.1347281$
- $p.value = 3.425363 \times 10^{-16}$

Question 2

Estimate an OLS regression in R that uses the Newton-Raphson algorithm (specifically BFGS, which is a quasi-Newton method), and show that you get the equivalent results to using `lm`. Use the code below to create your data.

```

1 set.seed (123)
2 data <- data.frame(x = runif(200, 1, 10))
3 data$y <- 0 + 2.75*data$x + rnorm(200, 0, 1.5)
4 lm_fit <- lm(y ~ x - 1, data = data)
5 coef(lm_fit)
6 ols_objective <- function(beta, y, x) {
7   sum((y - beta * x)^2)}
8 bfgs_fit <- optim(
9   par = 0,
10  fn = ols_objective,
11  y = data$y,
12  x = data$x,
13  method = "BFGS")
14 bfgs_fit$par
15 c(
16   lm_estimate = coef(lm_fit),
17   bfgs_estimate = bfgs_fit$par)
18 max(abs(
19   fitted(lm_fit) - data$x * bfgs_fit$par)))

```

The code first ensures that `data` refers to a numeric vector rather than a function, then creates the empirical cumulative distribution function `ECDF` for the observed data. A custom function, `ks normal test`, is defined to perform a Kolmogorov-Smirnov test against a normal distribution. Within this function, the KS statistic D is calculated as the maximum absolute difference between the empirical CDF and the theoretical normal CDF, and the corresponding p-value is approximated using the Kolmogorov distribution. The test is then applied to 1,000 randomly generated Cauchy values. The resulting large D and small p-value indicate a

substantial deviation from normality, leading to rejection of the null hypothesis that the data are normally distributed.

```
[1] 2.74764
> c(
+   lm_estimate = coef(lm_fit),
+   bfgs_estimate = bfgs_fit$par)
lm_estimate.x bfgs_estimate
2.74764      2.74764
> max(abs(
+ fitted(lm_fit) - data$x * bfgs_fit$par))
[1] 2.544591e-08
```